

Lecture 7

**Aggregate functions.
Window functions.**

Aggregate functions. Window functions.

PART I. Aggregate functions.

PART II. Window functions.

PART I. Aggregate functions.

Introduction to PostgreSQL aggregate functions

Aggregate functions perform a calculation on a set of rows and return a single row. PostgreSQL provides all standard SQL's aggregate functions as follows:

AVG() – return the average value.

COUNT() – return the number of values.

MAX() – return the maximum value.

MIN() – return the minimum value.

SUM() – return the sum of all or distinct values.

We often use the aggregate functions with the GROUP BY clause in the SELECT statement. In these cases, the GROUP BY clause divides the result set into groups of rows and the aggregate functions perform a calculation on each group e.g., maximum, minimum, average, etc.

You can use aggregate functions as expressions only in the following clauses:

SELECT clause.

HAVING clause.

PostgreSQL AVG Function

The **AVG()** function is one of the most commonly used aggregate functions in PostgreSQL. The AVG() function allows you to calculate the average value of a set.

The syntax of the AVG() function is as follows:

AVG(column)

You can use the AVG() function in the SELECT and HAVING clauses.

Example

```
SELECT AVG(amount)  
FROM payment;
```

| payment |
|--------------|
| * payment_id |
| customer_id |
| staff_id |
| rental_id |
| amount |
| payment_date |

```
avg
```

```
-----  
4.2006056453822965  
(1 row)
```

PostgreSQL AVG() function with DISTINCT operator

```
SELECT AVG(DISTINCT amount)::numeric(10,2)  
FROM payment;
```

| avg |
|---------|
| ----- |
| 6.14 |
| (1 row) |

PostgreSQL AVG function with SUM function

```
SELECT  
    AVG(amount)::numeric(10,2),  
    SUM(amount)::numeric(10,2)  
FROM  
    payment;
```

| avg | | sum |
|---------|--|----------|
| 4.20 | | 61312.04 |
| (1 row) | | |

PostgreSQL AVG() function with GROUP BY clause

```
SELECT
    customer_id,
    first_name,
    last_name,
    AVG(amount)::NUMERIC(10,2)
FROM
    payment
INNER JOIN customer USING(customer_id)
GROUP BY
    customer_id
ORDER BY
    customer_id;
```

| | customer_id integer | first_name character varying (45) | last_name character varying (45) | avg numeric (10,2) |
|----|------------------------|--------------------------------------|-------------------------------------|-----------------------|
| 1 | 1 | Mary | Smith | 3.82 |
| 2 | 2 | Patricia | Johnson | 4.76 |
| 3 | 3 | Linda | Williams | 5.45 |
| 4 | 4 | Barbara | Jones | 3.72 |
| 5 | 5 | Elizabeth | Brown | 3.85 |
| 6 | 6 | Jennifer | Davis | 3.39 |
| 7 | 7 | Maria | Miller | 4.67 |
| 8 | 8 | Susan | Wilson | 3.73 |
| 9 | 9 | Margaret | Moore | 3.94 |
| 10 | 10 | Dorothy | Taylor | 3.95 |
| 11 | 11 | Lisa | Anderson | 4.34 |
| 12 | 12 | Nancy | Thomas | 3.61 |
| 13 | 13 | Karen | Jackson | 4.88 |

PostgreSQL AVG() function with HAVING clause

```
SELECT
    customer_id,
    first_name,
    last_name,
    AVG(amount)::NUMERIC(10,2)
FROM
    payment
INNER JOIN customer USING(customer_id)
GROUP BY
    customer_id
HAVING
    AVG(amount) > 5
ORDER BY
    customer_id;
```

| | customer_id | first_name | last_name | avg |
|----|-------------|------------------------|------------------------|----------------|
| | integer | character varying (45) | character varying (45) | numeric (10,2) |
| 1 | 3 | Linda | Williams | 5.45 |
| 2 | 19 | Ruth | Martinez | 5.49 |
| 3 | 137 | Rhonda | Kennedy | 5.04 |
| 4 | 181 | Ana | Bradley | 5.08 |
| 5 | 187 | Brittany | Riley | 5.62 |
| 6 | 209 | Tonya | Chapman | 5.09 |
| 7 | 259 | Lena | Jensen | 5.16 |
| 8 | 272 | Kay | Caldwell | 5.07 |
| 9 | 285 | Miriam | Mckinney | 5.12 |
| 10 | 293 | Mae | Fletcher | 5.13 |
| 11 | 310 | Daniel | Cabral | 5.30 |
| 12 | 311 | Paul | Trout | 5.39 |

PostgreSQL AVG() function and NULL

```
CREATE TABLE t1 (
    id serial PRIMARY KEY,
    amount INTEGER
);
```

```
INSERT INTO t1 (amount)
VALUES
    (10),
    (NULL),
    (30);
```

```
SELECT AVG(amount)::numeric(10,2)
FROM t1;
```

| avg |
|---------|
| ----- |
| 20.00 |
| (1 row) |

PostgreSQL COUNT Function

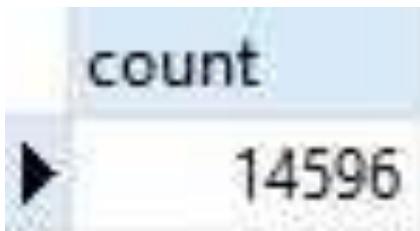
The **COUNT()** function is an aggregate function that allows you to get the number of rows that match a specific condition of a query.

```
SELECT COUNT(column)
FROM table_name
WHERE condition;
```

PostgreSQL COUNT() function examples

```
SELECT  
    COUNT(*)  
FROM  
    payment;
```

| payment |
|--------------|
| * payment_id |
| customer_id |
| staff_id |
| rental_id |
| amount |
| payment_date |



A screenshot of a PostgreSQL terminal window. The command `SELECT COUNT(*) FROM payment;` has been run, and the result is displayed in a table. The table has one row with the value `14596`. The column header is `count`.

| count |
|-------|
| 14596 |

PostgreSQL COUNT() with GROUP BY clause

```
SELECT  
    customer_id,  
    COUNT(customer_id)  
FROM  
    payment  
GROUP BY  
    customer_id;
```

| customer_id | count |
|-------------|-------|
| 251 | 25 |
| 106 | 21 |
| 120 | 30 |
| 285 | 23 |
| 264 | 25 |
| 497 | 27 |
| 452 | 29 |

PostgreSQL COUNT() with HAVING clause

```
SELECT  
    customer_id,  
    COUNT(customer_id)  
FROM  
    payment  
GROUP BY  
    customer_id  
HAVING  
    COUNT(customer_id) > 40;
```

| customer_id | count |
|-------------|-------|
| 526 | 42 |
| 148 | 45 |

PostgreSQL SUM Function

The PostgreSQL **SUM()** is an aggregate function that returns the sum of values or distinct values.

The syntax of the SUM() function is as follows:

`SUM(DISTINCT expression)`

PostgreSQL SUM() function examples

```
SELECT SUM(amount) AS total  
FROM payment  
WHERE customer_id = 2000;
```

| payment |
|--------------|
| * payment_id |
| customer_id |
| staff_id |
| rental_id |
| amount |
| payment_date |

```
total  
-----  
null  
(1 row)
```

Using PostgreSQL SUM() function with GROUP BY clause

```
SELECT  
    customer_id,  
    SUM(amount) AS total  
FROM  
    payment  
GROUP BY  
    customer_id  
ORDER BY total;
```

| | customer_id smallint | total numeric |
|----|-------------------------|------------------|
| 1 | 318 | 27.93 |
| 2 | 281 | 32.90 |
| 3 | 248 | 37.87 |
| 4 | 320 | 47.85 |
| 5 | 110 | 49.88 |
| 6 | 586 | 50.83 |
| 7 | 288 | 52.81 |
| 8 | 250 | 54.85 |
| 9 | 271 | 56.84 |
| 10 | 395 | 57.81 |
| 11 | 124 | 57.86 |

Using PostgreSQL SUM with expression

```
SELECT SUM(return_date - rental_date )  
FROM rental;
```

```
sum
```

```
-----  
71786 days 190098:21:00
```

```
(1 row)
```

| rental |
|--------------|
| * rental_id |
| rental_date |
| inventory_id |
| customer_id |
| return_date |
| staff_id |
| last_update |

PostgreSQL MAX Function

PostgreSQL **MAX** function is an aggregate function that returns the maximum value in a set of values.

The syntax of the MAX function is as follows:

`MAX(expression);`

You can use the MAX function not only in the SELECT clause but also in the WHERE and HAVING clauses.

PostgreSQL MAX function examples

```
SELECT MAX(amount)  
FROM payment;
```

| max |
|-------|
| 11.99 |

| payment |
|--------------|
| * payment_id |
| customer_id |
| staff_id |
| rental_id |
| amount |
| payment_date |

PostgreSQL MAX function in subquery

```
SELECT * FROM payment
```

```
WHERE amount = (
```

```
    SELECT MAX(amount)
```

```
    FROM payment
```

```
);
```

| payment_id | customer_id | staff_id | rental_id | amount | payment_date |
|------------|-------------|----------|-----------|--------|----------------------------|
| 20403 | 362 | 1 | 14759 | 11.99 | 2007-03-21 21:57:24.996577 |
| 22650 | 204 | 2 | 15415 | 11.99 | 2007-03-22 22:17:22.996577 |
| 23757 | 116 | 2 | 14763 | 11.99 | 2007-03-21 22:02:26.996577 |
| 24553 | 195 | 2 | 16040 | 11.99 | 2007-03-23 20:47:59.996577 |
| 24866 | 237 | 2 | 11479 | 11.99 | 2007-03-02 20:46:39.996577 |
| 28799 | 591 | 2 | 4383 | 11.99 | 2007-04-07 19:14:17.996577 |
| 28814 | 592 | 1 | 3973 | 11.99 | 2007-04-06 21:26:57.996577 |
| 29136 | 13 | 2 | 8831 | 11.99 | 2007-04-29 21:06:07.996577 |

PostgreSQL MIN Function

PostgreSQL **MIN()** function an aggregate function that returns the minimum value in a set of values.

The syntax of the MIN() function is as follows:

SELECT

 MIN(expression)

FROM

 table_expression

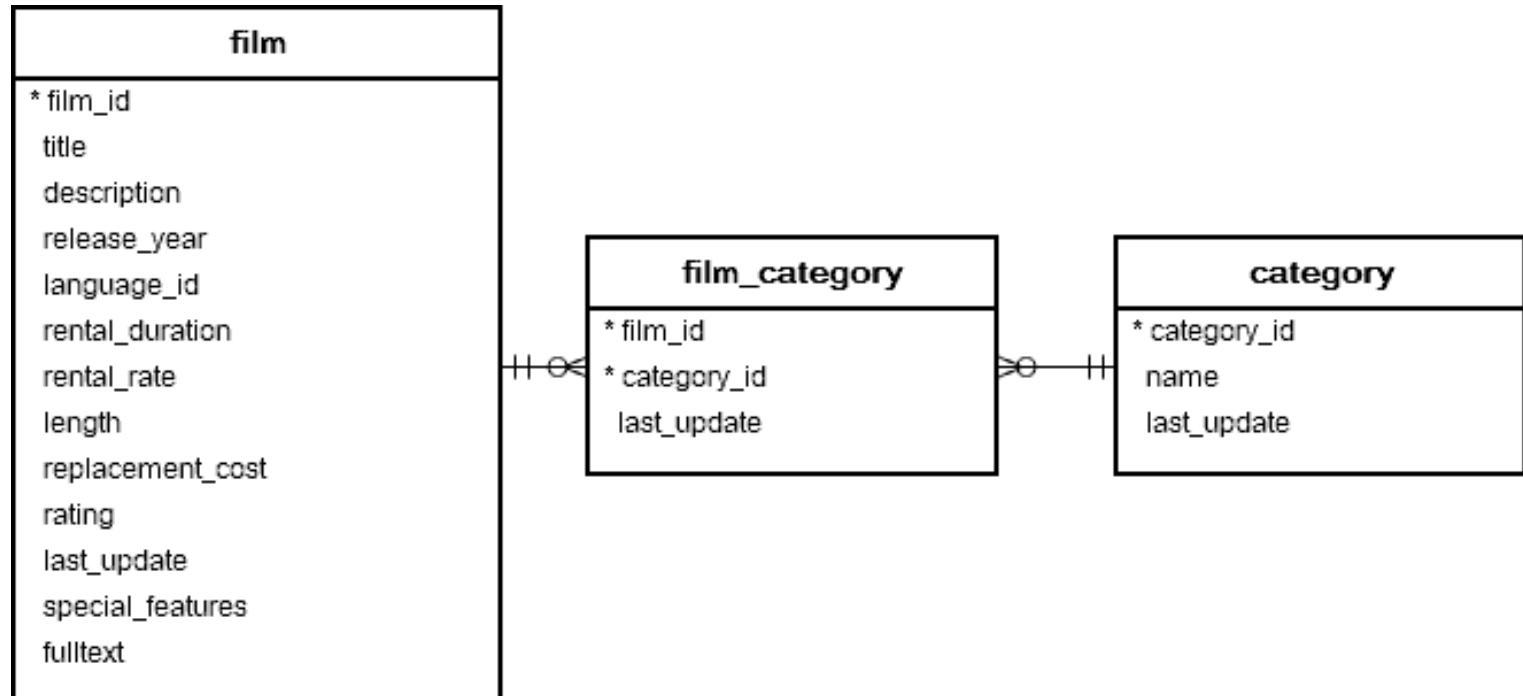
...;

Unlike the AVG(), COUNT() and SUM() functions, the DISTINCT option does not have any effects on the MIN() function.

PostgreSQL MIN() function examples

```
SELECT  
    MIN (rental_rate)  
FROM  
    film;
```

| | min |
|---|---------|
| | numeric |
| 1 | 0.99 |



Using PostgreSQL MIN function with GROUP BY clause

```
SELECT
    name category,
    MIN(replacement_cost) replacement_cost
FROM category
INNER JOIN film_category USING (category_id)
INNER JOIN film USING (film_id)
GROUP BY name
ORDER BY name;
```

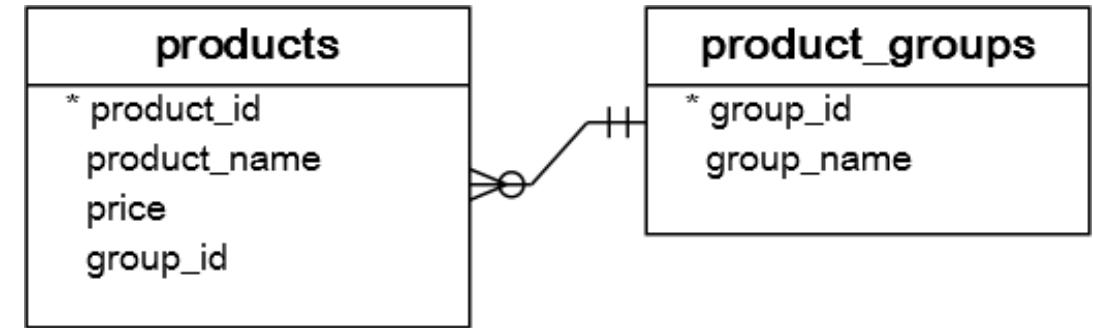
| | category character varying (25) | replacement_cost numeric |
|----|------------------------------------|-----------------------------|
| 1 | Action | 9.99 |
| 2 | Animation | 9.99 |
| 3 | Children | 9.99 |
| 4 | Classics | 10.99 |
| 5 | Comedy | 9.99 |
| 6 | Documentary | 9.99 |
| 7 | Drama | 9.99 |
| 8 | Family | 9.99 |
| 9 | Foreign | 9.99 |
| 10 | Games | 9.99 |
| 11 | Horror | 10.99 |
| 12 | Music | 10.99 |
| 13 | New | 9.99 |
| 14 | Sci-Fi | 9.99 |
| 15 | Sports | 9.99 |
| 16 | Travel | 9.99 |

PART II. Window functions.

PostgreSQL Window Functions

```
CREATE TABLE product_groups (
    group_id serial PRIMARY KEY,
    group_name VARCHAR (255) NOT NULL
);
```

```
CREATE TABLE products (
    product_id serial PRIMARY KEY,
    product_name VARCHAR (255) NOT NULL,
    price DECIMAL (11, 2),
    group_id INT NOT NULL,
    FOREIGN KEY (group_id) REFERENCES product_groups (group_id)
);
```



PostgreSQL Window Functions

```
INSERT INTO product_groups (group_name)
```

```
VALUES
```

```
    ('Smartphone'),  
    ('Laptop'),  
    ('Tablet');
```

```
INSERT INTO products (product_name, group_id, price)
```

```
VALUES
```

```
    ('Microsoft Lumia', 1, 200),  
    ('HTC One', 1, 400),  
    ('Nexus', 1, 500),  
    ('iPhone', 1, 900),  
    ('HP Elite', 2, 1200),  
    ('Lenovo Thinkpad', 2, 700),  
    ('Sony VAIO', 2, 700),  
    ('Dell Vostro', 2, 800),  
    ('iPad', 3, 700),  
    ('Kindle Fire', 3, 150),  
    ('Samsung Galaxy Tab', 3, 200);
```

Introduction to PostgreSQL window functions

```
SELECT
    group_name,
    AVG(price)
FROM
    products
INNER JOIN product_groups USING (group_id)
GROUP BY
    group_name;
```

| group_name | avg |
|------------|-----|
| Smartphone | 500 |
| Tablet | 350 |
| Laptop | 850 |

Example

```
SELECT
    product_name,
    price,
    group_name,
    AVG(price) OVER (
        PARTITION BY group_name
    )
```

```
FROM
    products
INNER JOIN
    product_groups USING (group_id);
```

| product_name | price | group_name | avg |
|--------------------|-------|------------|-----|
| HP Elite | 1200 | Laptop | 850 |
| Lenovo Thinkpad | 700 | Laptop | 850 |
| Sony VAIO | 700 | Laptop | 850 |
| Dell Vostro | 800 | Laptop | 850 |
| Microsoft Lumia | 200 | Smartphone | 500 |
| HTC One | 400 | Smartphone | 500 |
| Nexus | 500 | Smartphone | 500 |
| iPhone | 900 | Smartphone | 500 |
| iPad | 700 | Tablet | 350 |
| Kindle Fire | 150 | Tablet | 350 |
| Samsung Galaxy Tab | 200 | Tablet | 350 |

PostgreSQL Window Function

A **window function** performs a calculation across a set of table rows that are somehow related to the current row.

PostgreSQL Window Function Syntax:

```
window_function(arg1, arg2,...) OVER (  
    [PARTITION BY partition_expression]  
    [ORDER BY sort_expression [ASC | DESC] [NULLS {FIRST | LAST }]])
```

Window functions

| Name | Description |
|--------------|---|
| CUME_DIST | Return the relative rank of the current row. |
| DENSE_RANK | Rank the current row within its partition without gaps. |
| FIRST_VALUE | Return a value evaluated against the first row within its partition. |
| LAG | Return a value evaluated at the row that is at a specified physical offset row before the current row within the partition. |
| LAST_VALUE | Return a value evaluated against the last row within its partition. |
| LEAD | Return a value evaluated at the row that is offset rows after the current row within the partition. |
| NTILE | Divide rows in a partition as equally as possible and assign each row an integer starting from 1 to the argument value. |
| NTH_VALUE | Return a value evaluated against the nth row in an ordered partition. |
| PERCENT_RANK | Return the relative rank of the current row (rank-1) / (total rows – 1) |
| RANK | Rank the current row within its partition with gaps. |
| ROW_NUMBER | Number the current row within its partition starting from 1. |

The ROW_NUMBER() function

The **ROW_NUMBER()** function assigns a sequential number to each row in each partition.
See the following query:

```
SELECT
    product_name,
    group_name,
    price,
    ROW_NUMBER () OVER (
        PARTITION BY group_name
        ORDER BY
            price
    )
FROM
    products
INNER JOIN product_groups USING (group_id);
```

| product_name | group_name | price | row_number |
|--------------------|------------|-------|------------|
| Sony VAIO | Laptop | 700 | 1 |
| Lenovo Thinkpad | Laptop | 700 | 2 |
| Dell Vostro | Laptop | 800 | 3 |
| HP Elite | Laptop | 1200 | 4 |
| Microsoft Lumia | Smartphone | 200 | 1 |
| HTC One | Smartphone | 400 | 2 |
| Nexus | Smartphone | 500 | 3 |
| iPhone | Smartphone | 900 | 4 |
| Kindle Fire | Tablet | 150 | 1 |
| Samsung Galaxy Tab | Tablet | 200 | 2 |
| iPad | Tablet | 700 | 3 |

The Rank() function

The **RANK()** function assigns ranking within an ordered partition. If rows have the same values, the **RANK()** function assigns the same rank, with the next ranking(s) skipped.

SELECT

```
    product_name,  
    group_name,  
    price,  
    RANK () OVER (  
        PARTITION BY group_name  
        ORDER BY  
            price  
    )  
FROM  
    products  
INNER JOIN product_groups USING (group_id);
```

| | product_name | group_name | price | rank |
|---|--------------------|------------|-------|------|
| ▶ | Sony VAIO | Laptop | 700 | 1 |
| | Lenovo Thinkpad | Laptop | 700 | 1 |
| | Dell Vostro | Laptop | 800 | 3 |
| | HP Elite | Laptop | 1200 | 4 |
| | Microsoft Lumia | Smartphone | 200 | 1 |
| | HTC One | Smartphone | 400 | 2 |
| | Nexus | Smartphone | 500 | 3 |
| | iPhone | Smartphone | 900 | 4 |
| | Kindle Fire | Tablet | 150 | 1 |
| | Samsung Galaxy Tab | Tablet | 200 | 2 |
| | iPad | Tablet | 700 | 3 |

The Dense_rank() function

Similar to the RANK() function, the **DENSE_RANK()** function assigns a rank to each row within an ordered partition, but the ranks have no gap. In other words, the same ranks are assigned to multiple rows and no ranks are skipped.

```
SELECT
    product_name,
    group_name,
    price,
    DENSE_RANK () OVER (
        PARTITION BY group_name
        ORDER BY
            price
    )
FROM
    products
INNER JOIN product_groups USING (group_id);
```

| product_name | group_name | price | dense_rank |
|--------------------|------------|-------|------------|
| Sony VAIO | Laptop | 700 | 1 |
| Lenovo Thinkpad | Laptop | 700 | 1 |
| Dell Vostro | Laptop | 800 | 2 |
| HP Elite | Laptop | 1200 | 3 |
| Microsoft Lumia | Smartphone | 200 | 1 |
| HTC One | Smartphone | 400 | 2 |
| Nexus | Smartphone | 500 | 3 |
| iPhone | Smartphone | 900 | 4 |
| Kindle Fire | Tablet | 150 | 1 |
| Samsung Galaxy Tab | Tablet | 200 | 2 |
| iPad | Tablet | 700 | 3 |

The FIRST_VALUE() function

The **FIRST_VALUE()** function returns a value evaluated against the first row within its partition.

```
SELECT
    product_name,
    group_name,
    price,
    FIRST_VALUE (price) OVER (
        PARTITION BY group_name
        ORDER BY
            price
    ) AS lowest_price_per_group
FROM
    products
INNER JOIN product_groups USING (group_id);
```

| product_name | group_name | price | lowest_price_per_group |
|--------------------|------------|-------|------------------------|
| Sony VAIO | Laptop | 700 | 700 |
| Lenovo Thinkpad | Laptop | 700 | 700 |
| Dell Vostro | Laptop | 800 | 700 |
| HP Elite | Laptop | 1200 | 700 |
| Microsoft Lumia | Smartphone | 200 | 200 |
| HTC One | Smartphone | 400 | 200 |
| Nexus | Smartphone | 500 | 200 |
| iPhone | Smartphone | 900 | 200 |
| Kindle Fire | Tablet | 150 | 150 |
| Samsung Galaxy Tab | Tablet | 200 | 150 |
| iPad | Tablet | 700 | 150 |

The LAST_VALUE() function

The **LAST_VALUE()** function returns a value evaluated against the last row in its partition.

SELECT

```
product_name,  
group_name,  
price,  
LAST_VALUE (price) OVER (  
    PARTITION BY group_name  
    ORDER BY  
        price  
) AS highest_price_per_group
```

FROM

products

INNER JOIN product_groups USING (group_id);

| product_name | group_name | price | highest_price_per_group |
|--------------------|------------|-------|-------------------------|
| Sony VAIO | Laptop | 700 | 1200 |
| Lenovo Thinkpad | Laptop | 700 | 1200 |
| Dell Vostro | Laptop | 800 | 1200 |
| HP Elite | Laptop | 1200 | 1200 |
| Microsoft Lumia | Smartphone | 200 | 900 |
| HTC One | Smartphone | 400 | 900 |
| Nexus | Smartphone | 500 | 900 |
| iPhone | Smartphone | 900 | 900 |
| Kindle Fire | Tablet | 150 | 700 |
| Samsung Galaxy Tab | Tablet | 200 | 700 |
| iPad | Tablet | 700 | 700 |

The LAG and LEAD functions

The **LAG()** function has the ability to access data from the previous row, while the **LEAD()** function can access data from the next row.

Both LAG() and LEAD() functions have the same syntax as follows:

```
LAG (expression [,offset] [,default]) over_clause;  
LEAD (expression [,offset] [,default]) over_clause;
```

In this syntax:

- expression – a column or expression to compute the returned value.
- offset – the number of rows preceding (LAG)/ following (LEAD) the current row. It defaults to 1.
- default – the default returned value if the offset goes beyond the scope of the window. The default is NULL if you skip it.

The LAG() function example

SELECT

```
product_name,  
group_name,  
price,  
LAG(price, 1) OVER (  
    PARTITION BY group_name  
    ORDER BY  
        price  
) AS prev_price,  
price - LAG(price, 1) OVER (  
    PARTITION BY group_name  
    ORDER BY  
        price  
) AS cur_prev_diff
```

FROM

products

INNER JOIN product_groups USING (group_id);

| product_name | group_name | price | prev_price | cur_prev_diff |
|--------------------|------------|-------|------------|---------------|
| Sony VAIO | Laptop | 700 | (Null) | (Null) |
| Lenovo Thinkpad | Laptop | 700 | 700 | 0 |
| Dell Vostro | Laptop | 800 | 700 | 100 |
| HP Elite | Laptop | 1200 | 800 | 400 |
| Microsoft Lumia | Smartphone | 200 | (Null) | (Null) |
| HTC One | Smartphone | 400 | 200 | 200 |
| Nexus | Smartphone | 500 | 400 | 100 |
| iPhone | Smartphone | 900 | 500 | 400 |
| Kindle Fire | Tablet | 150 | (Null) | (Null) |
| Samsung Galaxy Tab | Tablet | 200 | 150 | 50 |
| iPad | Tablet | 700 | 200 | 500 |

The LEAD() function example

SELECT

product_name, group_name, price,
LEAD (price, 1) OVER (

PARTITION BY group_name ORDER BY product_name

) AS next_price,

price - LEAD (price, 1) OVER (PARTITION BY group_name ORDER BY product_name

) AS cur_next_diff

FROM

products

INNER JOIN product_groups USING (group_id);

| product_name | group_name | price | next_price | cur_next_diff |
|--------------------|------------|-------|------------|---------------|
| Sony VAIO | Laptop | 700 | 700 | 0 |
| Lenovo Thinkpad | Laptop | 700 | 800 | -100 |
| Dell Vostro | Laptop | 800 | 1200 | -400 |
| HP Elite | Laptop | 1200 | (Null) | (Null) |
| Microsoft Lumia | Smartphone | 200 | 400 | -200 |
| HTC One | Smartphone | 400 | 500 | -100 |
| Nexus | Smartphone | 500 | 900 | -400 |
| iPhone | Smartphone | 900 | (Null) | (Null) |
| Kindle Fire | Tablet | 150 | 200 | -50 |
| Samsung Galaxy Tab | Tablet | 200 | 700 | -500 |
| iPad | Tablet | 700 | (Null) | (Null) |